



INSTITUTO SUPERIOR TÉCNICO

DEPARTAMENTO DE ENGENHARIA INFORMÁTICA

FORENSICS CYBER-SECURITY

MEIC, METI

Tutorial III

Introduction to File System Forensics

2020/2021

nuno.m.santos@tecnico.ulisboa.pt

Introduction

This guide aims to provide a hands-on introduction to file system forensics and to introduce you to the usage of well-known file system analysis tools such as The Sleuth Kit (TSK), as well as Scalpel and Foremost, two file carving tools which may help to collect evidence in the case file metadata is not available. The forensic images required for these exercises can be downloaded from the course website.

Before we begin, make sure to prepare the environment for forensic analysis. Please refer to the Kali Linux forensic testbed that you have set up in Tutorial I and reproduce it before proceeding with the remaining of this guide. We also suggest you to review the slides of theory classes 9 and 10 which provide you the technical background that will allow you to understand the steps presented herein.

1 Deleted File Identification and Recovery (Ext2)

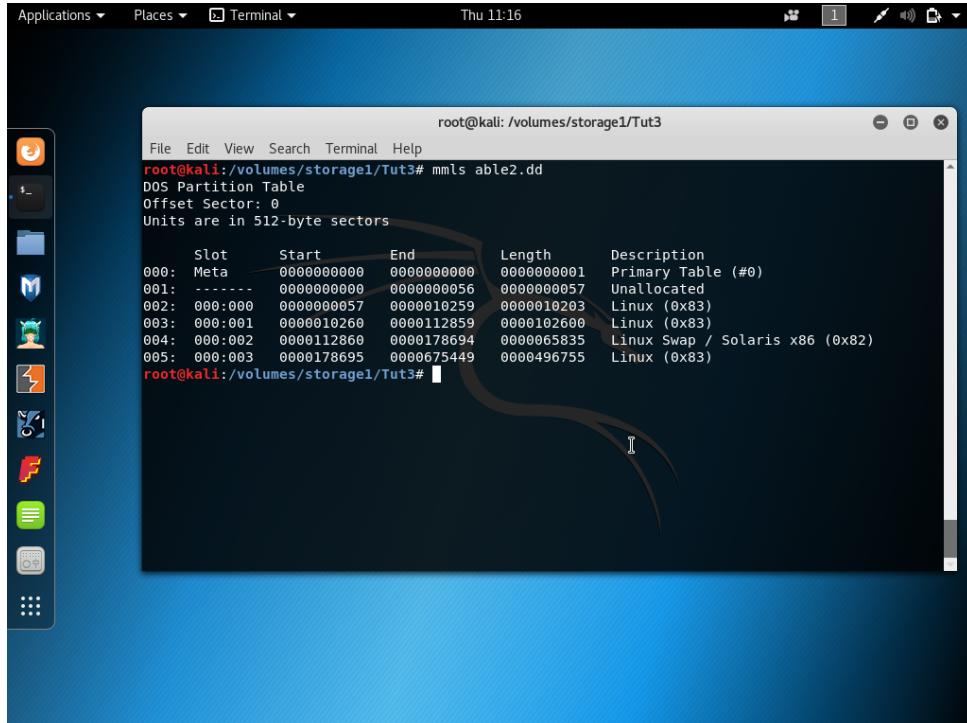
The goal of this exercise is to train you to the identification and recovery of deleted files from a forensic disk image containing an Ext2 file system. In this particular case, the file can be recovered based on leftover unallocated metadatada that can be retrieved using The Sleuth Kit (TSK). TSK is a collection of command line tools and a C library that allows you to analyze disk images and recover files from them. We will start by getting familiar with a couple of file system analysis tools provided by TSK, like `fsstat`, `f1s`, or `mmls`, running them against a disk image.

Before starting your analysis, download the `able2` disk image into some folder in the Kali forensic environment (e.g., `/volumes/storage1/Tut3`) by executing the following command:

```
root@kali:/volumes/storage1/Tut3# wget http://turbina.gsd.inesc-id.pt/csf2021/able2.tar.gz
```

1.1 List the partition table

The first tool we are going to use, `mmls`, provides access to the partition table within an image, and gives the partition offsets in sector units. The output of `mmls able2.dd` can be observed in Figure 1.



The screenshot shows a terminal window on a Kali Linux desktop. The window title is "root@kali: /volumes/storage1/Tut3#". The command entered is "mmls able2.dd". The output is as follows:

```
File Edit View Search Terminal Help
root@kali:/volumes/storage1/Tut3# mmls able2.dd
DOS Partition Table
Offset Sector: 0
Units are in 512-byte sectors

      Slot      Start        End      Length     Description
 000: Meta  0000000000  0000000000  0000000001 Primary Table (#0)
 001: ----- 0000000000  0000000056  0000000057 Unallocated
 002: 000:000 0000000057  0000010259  0000010203 Linux (0x83)
 003: 000:001 0000010260  0000112859  0000102600 Linux (0x83)
 004: 000:002 00000112860  00000178694  00000065835 Linux Swap / Solaris x86 (0x82)
 005: 000:003 00000178695  00000675449  0000496755 Linux (0x83)

root@kali:/volumes/storage1/Tut3#
```

Figure 1: Output of `mmls able2.dd`

In this analysis, suppose that the information we are looking for is located on the root file system of our image. The root (/) file system is located on the second partition. Looking at our `mmls` output, we can see that that partition starts at sector 10260 (numbered 03 in the `mmls` output, or slot 000:001).

1.2 Gather file system information

Next, we need to learn what's inside our target partition. The `fsstat` command provides specific information about the file system in a volume. We can run TSK's `fsstat` command with `-o 10260` to gather file system information at that offset (Figure 2). This specifies that we want information residing on the partition that starts at sector offset 10260.

Figure 2: Output of `fsstat -o 10260 able2.dd | less`

1.3 List file names and directories

We can get more information using the `fls` command, which lists the file names and directories contained in a file system, or in a given directory. If you run the `fls` command with only the `fls` option to specify the file system, then by default it will run on the file system's root directory (see Figure 3). This is inode 2 on an EXT file system and MFT entry 5 on an NTFS file system. In other words, on an EXT file system, any of the below commands will yield the same output (Fig 3):

```
root@kali:/volumes/storage1/Tut3# fls -o 10260 able2.dd
```

or

```
root@kali:/volumes/storage1/Tut3# fls -o 10260 able2.dd 2
```

There are several points we want to take note of before we continue. Let's take a few lines of output and describe what the tool is telling us. Have a look at the last three lines from the above `fls` command.

```
...
r/r 1042: .bash_history
d/d 11105: .001
d7d 12881 $OrphanFiles
```

```
root@kali: /volumes/storage1/Tut3#
File Edit View Search Terminal Help
004: 000:002 0000112860 0000178694 0000065835 Linux Swap / Solaris x86
(0x82)
005: 000:003 0000178695 0000675449 0000496755 Linux (0x83)
root@kali:/volumes/storage1/Tut3# fsstat -o 10260 able2.dd | less
root@kali:/volumes/storage1/Tut3# fls -o 10260 able2.dd
d/d 11: lost+found
d/d 3681: boot
d/d 7361: usr
d/d 3682: proc
d/d 7362: var
d/d 5521: tmp
d/d 7363: dev
d/d 9201: etc
d/d 1843: bin
d/d 1844: home
d/d 7368: lib
d/d 7369: mnt
d/d 7370: opt
d/d 1848: root
d/d 1849: sbin
r/r 1042: .bash_history
d/d 11105: .001
d/d 12881: $OrphanFiles
root@kali:/volumes/storage1/Tut3#
```

Figure 3: Output of `fls -o 10260 able2.dd`

Each line of output starts with two characters separated by a slash. This field indicates the file type as described by the file's directory entry, and the file's metadata (in this case, the inode because we are looking at an EXT file system). For example, the first file listed in the snippet above, `.bash_history`, is identified as a regular file in both the file's directory and inode entry. This is noted by the `r/r` designation. Conversely, the following two entries (`.001` and `$OrphanFiles`) are identified as directories. The next field is the metadata entry number (inode, MFT entry, etc.) followed by the filename. In the case of the file `.bash_history` the inode is listed as `1042`.

Note that the last line of the output, `$OrphanFiles` is a virtual folder created by TSK and assigned a virtual inode. This folder contains virtual file entries that represent unallocated metadata entries where there are no corresponding file names. These are commonly referred to as “orphan files”, which can be accessed by specifying the metadata address, but not through any file name path.

```
root@kali: /volumes/storage1/Tut3#
File Edit View Search Terminal Help
root@kali:/volumes/storage1/Tut3# fls -o 10260 able2.dd 11105
r/r 2130: lolit_pics.tar.gz
r/r 11107: lolitaZ1
r/r 11108: lolitaZ10
r/r 11109: lolitaZ11
r/r 11110: lolitaZ12
r/r 11111: lolitaZ13
r/r 11112: lolitaZ2
r/r 11113: lolitaZ3
r/r 11114: lolitaZ4
r/r 11115: lolitaZ5
r/r 11116: lolitaZ6
r/r 11117: lolitaZ7
r/r 11118: lolitaZ8
r/r 11119: lolitaZ9
root@kali:/volumes/storage1/Tut3#
```

Figure 4: Output of `fls -o 10260 able2.dd 11105`

In Figure 4, we continue to run `f1s` on directory entries to dig deeper into the file system structure (we could also have used `-r` for a recursive listing). By passing the metadata entry number of a directory, we can view its contents. For example, have a look at the `.001` directory in the listing above. This is an unusual directory and would cause some suspicion. It is hidden (starts with a `.`), and no such directory is common in the root of the file system. To see the contents of the `.001` directory, we would pass its inode to `f1s` as it can be observed in Figure 4.

1.4 Uncover deleted files

Mostly important for this exercise, `f1s` can also be useful for uncovering deleted files. For example, if we wanted to exclusively check deleted entries that are listed as files (rather than directories), and we want the listing to be recursive, we could use run `f1s -o 10260 able2.dd -Frd able2.dd`. This command runs against the partition in `able2.dd` starting at sector offset 10260 (`-o 10260`), showing only file entries (`-F`), descending into directories recursively (`-r`), and displaying deleted entries (`-d`). The output of the above command can be observed in Figure 5.

```
root@kali: /volumes/storage1/Tut3# f1s -o 10260 able2.dd -Frd able2.dd
r/r * 11120(realloc): var/lib/slocate/slocate.db.tmp
r/r * 10063: var/log/xferlog.5
r/r * 10063: var/lock/makewhatistis.lock
r/r * 6613: var/run/shutdown.pid
r/r * 1046: var/tmp/rpm-tmp.64655
r/r * 6609(realloc): var/catman/cat1/rdate.1.gz
r/r * 6613: var/catman/cat1/rdate.1.gz
r/r * 6616: tmp/logrot2V6Q1J
r/r * 2139: dev/ttyZ0/lrkN.tgz
d/r * 10071(realloc): dev/ttyZ0/lrk3
r/r * 6572(realloc): etc/X11/fs/config-
l/r * 1041(realloc): etc/rc.d/rc0.d/K83ypbind
l/r * 1042(realloc): etc/rc.d/rc1.d/K83ypbind
l/r * 6583(realloc): etc/rc.d/rc2.d/K83ypbind
l/r * 6584(realloc): etc/rc.d/rc4.d/K83ypbind
l/r * 1044: etc/rc.d/rc5.d/K83ypbind
l/r * 6585(realloc): etc/rc.d/rc6.d/K83ypbind
r/r * 1044: etc/rc.d/rc.firewall-
r/r * 6544(realloc): etc/pam.d/passwd-
r/r * 10055(realloc): etc/mtab.tmp
r/r * 10047(realloc): etc/mtab-
r/- * 0: etc/.inetd.conf.swx
r/r * 2138(realloc): root/lolit_pics.tar.gz
```

Figure 5: Output of `f1s -o 10260 able2.dd -Frd able2.dd`

Notice that all of the files listed have an asterisk (*) before the inode. This indicates the file is deleted, which we expect in the above output since we specified the `-d` option to `f1s`. We are then presented with the metadata entry number (inode, MFT entry, etc.) followed by the filename.

Have a look at the line of output for inode number 2138 (`root/lolit_pics.tar.gz`). The inode is followed by `realloc`. Keep in mind that `f1s` describes the file name layer. The `realloc` means that the file name listed is marked as unallocated, even though the metadata entry (2138) is marked as allocated. Thus, the inode from our deleted file may have been “reallocated” to a new file.

In the case of inode 2138, it looks as though the `realloc` was caused by the file being moved to the directory `.001` (see the `f1s` listing of `.001` on the previous page – inode 11105). This causes it to be deleted from its current directory entry (`root/lolit_pics.tar.gz`) and a new file name created (`.001/lolit_pics.tar.gz`). The inode and the data blocks that it points to remain unchanged and in “allocated status”, but it has been “reallocated” to the new name.

1.5 Find all file names associated with a particular inode

Let's continue our analysis by taking advantage of metadata (inode) layer tools included in TSK. In a Linux EXT type file system, an inode has a unique number and is assigned to a file. The number corresponds to the inode table, allocated when a partition is formatted. The inode contains all the metadata available for a file, including the modified/accessible/changed times and a list of all the data blocks allocated to that file. If you look at the output of our last `f1s` command, you will see a deleted file called `lrkn.tgz` located in the `/root` directory:

```
...  
r/r * 2139:    root/lrkn.tgz  
...
```

The inode displayed by `f1s` for this file is 2139. This inode also points to another deleted file in `/dev` (same file, different location). We can find all the file names associated with a particular metadata entry by using the `ffind -o 10260 -a able2.dd 2139` command (Figure 6).

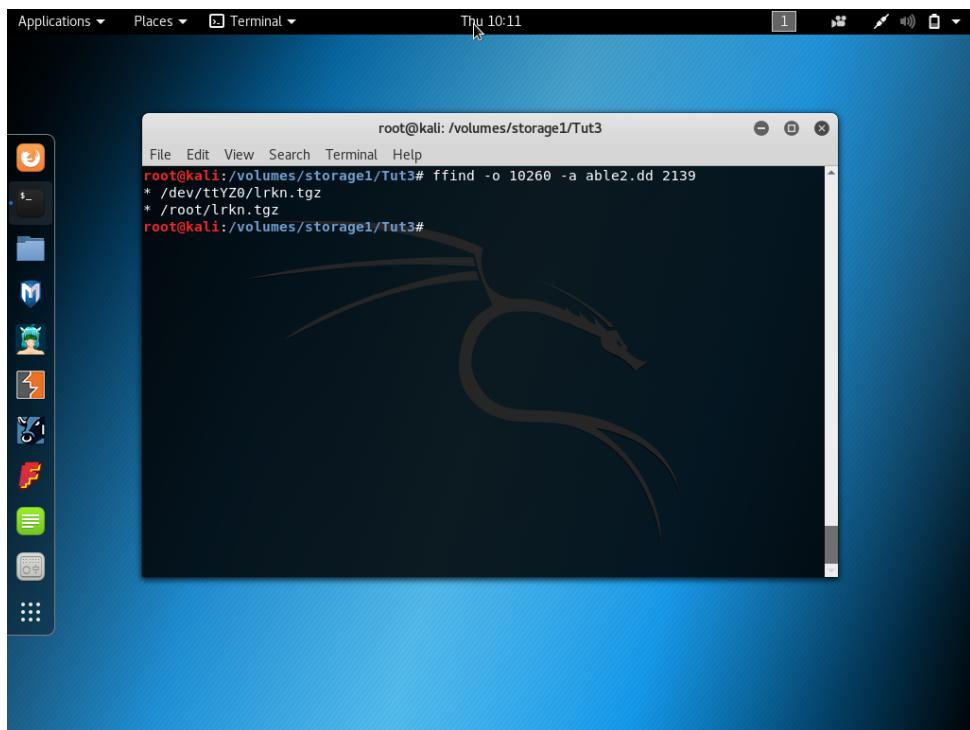


Figure 6: Output of `ffind -o 10260 -a able2.dd 2139`

Here we see that there are two file names associated with inode 2139, and both are deleted, as noted again by the asterisk (the `-a` ensures that we get all the inode associations).

1.6 Gather information about an inode

Continuing on, we are going to use `istat`. Remember that `fsstat` took a file system as an argument and reported statistics about that file system. `istat` does the same thing; only it works on a specified inode or metadata entry. In NTFS, this would be an MFT entry, for example.

We may execute `istat -o 10260 able2.dd 2139` to gather information about inode 2139 (Figure 7). This command reads the inode statistics on the file system located in the `able2.dd` image in the partition at sector offset 10260 (`-o 10260`), from inode 2139 found in our `f1s` command. There is a large amount of output here, showing all the inode information and the file system blocks ("Direct Blocks") that contain all of the file's data.

```
root@kali:/volumes/storage1/Tut3
File Edit View Search Terminal Help
inode: 2139
Not Allocated
Group: 1
Generation Id: 3534950564
uid / gid: 0 / 0
mode: rrw-r--r-
size: 3639016
num of links: 0

Inode Times:
Accessed: 2003-08-10 04:18:38 (UTC)
File Modified: 2003-08-10 04:08:32 (UTC)
Inode Modified: 2003-08-10 04:29:58 (UTC)
Deleted: 2003-08-10 04:29:58 (UTC)

Direct Blocks:
22811 22812 22813 22814 22815 22816 22817 22818
22819 22820 22821 22822 22824 22825 22826 22827
22828 22829 22830 22831 22832 22833 22834 22835
22836 22837 22838 22839 22840 22841 22842 22843
22844 22845 22846 22847 22848 22849 22850 22851
22852 22853 22854 22855 22856 22857 22858 22859
22860 22861 22862 22863 22864 22865 22866 22867
```

Figure 7: Output of `istat -o 10260 able2.dd 2139 | less`

1.7 Extract and inspect the deleted file based on its former inode

We now have the name of a deleted file of interest (from `f1s`) and the inode information, including where the data is stored (from `iStat`). Next we are going to use the `icat` command from TSK to grab the actual data contained in the data blocks referenced from the inode. `icat` also takes the inode as an argument and reads the content of the data blocks that are assigned to that inode, sending it to standard output. Remember, this is a deleted file that we are recovering here. We are going to send the contents of the data blocks assigned to inode 2139 to a file for closer examination:

```
root@kali:/volumes/storage1/Tut3# icat -o 10260 able2.dd 2139 > lrkn.tgz
.2139
```

This command runs the `icat` tool on the file system in our `able2.dd` image at sector offset 10260 (`-o 10260`) and streams the contents of the data blocks associated with inode 2139 to the file `lrkn.tgz.2139`. Now that we have what we hope is a recovered file, we can inspect the recovered data with the `file` command (Figure 8).

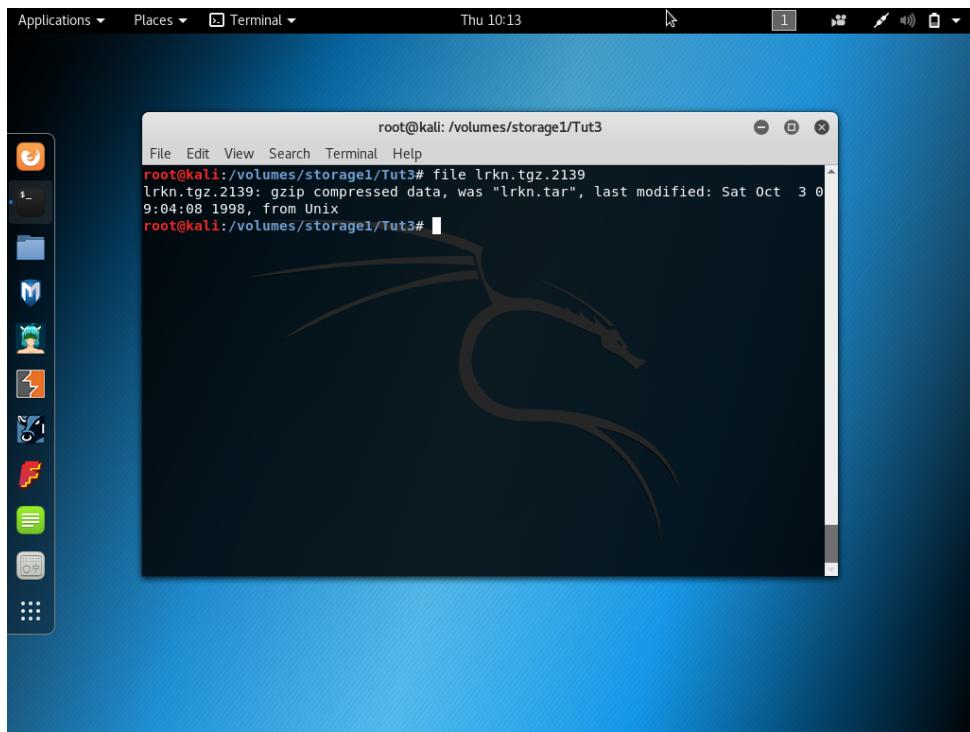


Figure 8: Output of file lrkn.tgz.2139

Have a look at the contents of the recovered archive by executing `tar tzvf lrkn.tgz.2139`. Recall that the `t` option to the tar command lists the contents of the archive (Figure 9).

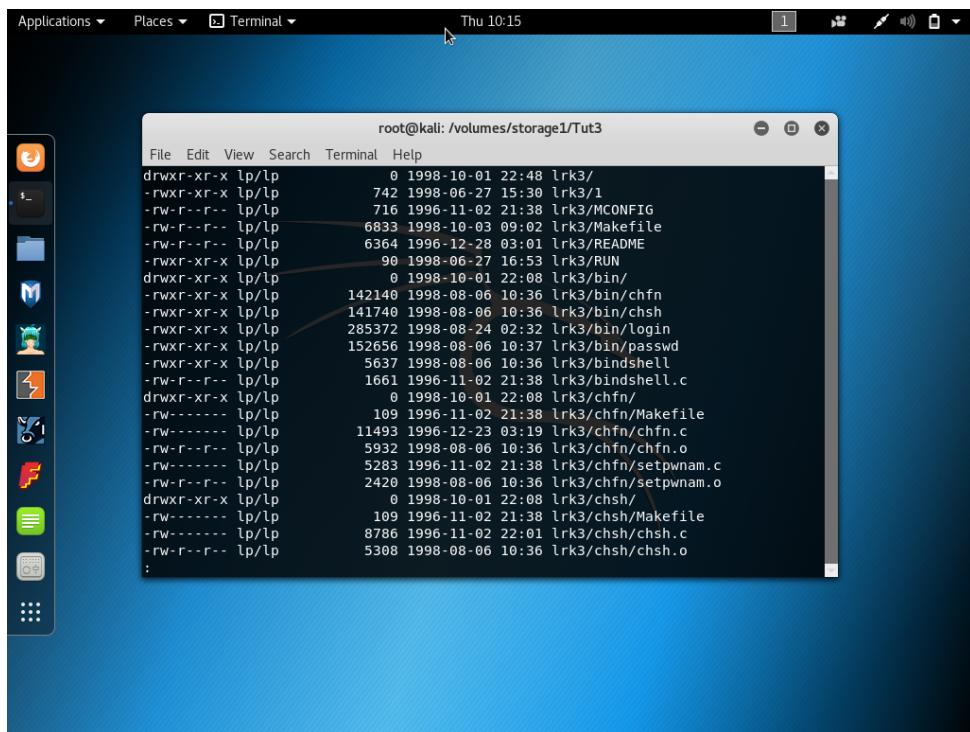
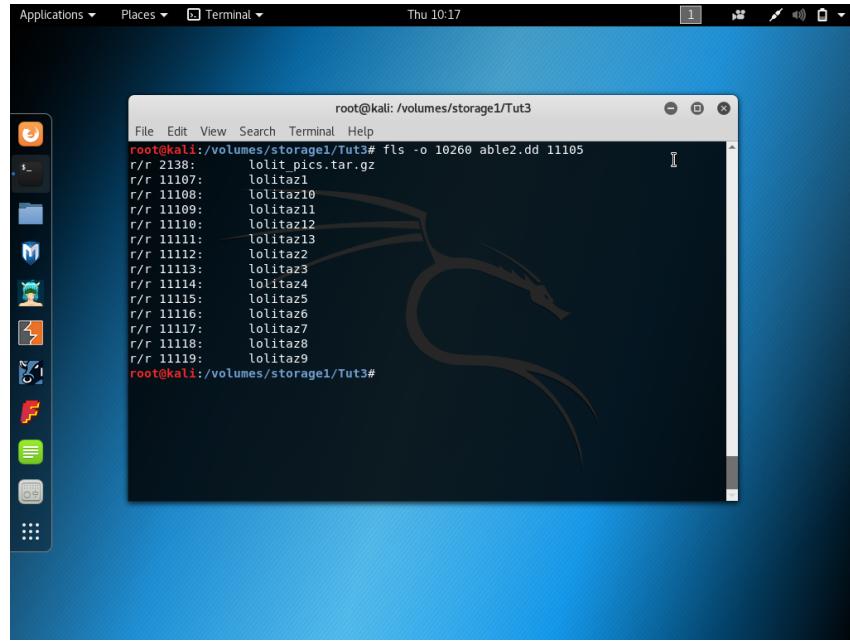


Figure 9: Output of tar tzvf lrkn.tgz.2139 | less

We can notice that there is a README file included in the archive. Rather than extracting the entire contents of the archive, we will first extract and inspect the README file using the command `tar xzvfO lrkn.tgz.2139 lrk3/README > lrkn.2139 README`. If you read the file, you will find out that we have uncovered a “rootkit”, full of programs used to hide a hacker’s activity.

1.8 More on icat

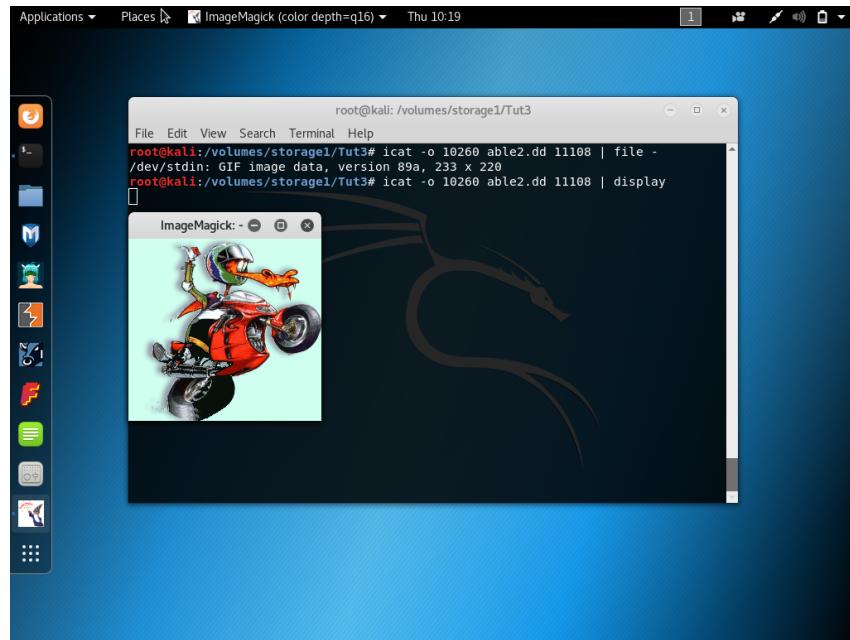
Let's now look at a different type of files recovered by `icat`. Recall our previous directory listing of the `.001` directory at inode 11105 (`fls -o 10260 able2.dd 11105`), depicted in Figure 10.



```
root@kali: /volumes/storage1/Tut3# fls -o 10260 able2.dd 11105
r/r 2138: lolit_pics.tar.gz
r/r 11107: lolita1
r/r 11108: lolita10
r/r 11109: lolita11
r/r 11110: lolita12
r/r 11111: lolita13
r/r 11112: lolita22
r/r 11113: lolita3
r/r 11114: lolita4
r/r 11115: lolita5
r/r 11116: lolita6
r/r 11117: lolita7
r/r 11118: lolita8
r/r 11119: lolita9
root@kali:/volumes/storage1/Tut3#
```

Figure 10: Output of `fls -o 10260 able2.dd 11105`

We can determine the contents of the (allocated) file with inode 11108 by using `icat` to stream the inode's data blocks through a pipe to `file` (`icat -o 10260 able2.dd 11108 | file -`). The output shows that we are dealing with a picture file. You may use the `display` command to show its contents by executing `icat -o 10260 able2.dd 11108 | display` (Figure 11).



```
root@kali: /volumes/storage1/Tut3# icat -o 10260 able2.dd 11108 | file -
/dev/stdin: GIF image data, version 89a, 233 x 220
root@kali:/volumes/storage1/Tut3# icat -o 10260 able2.dd 11108 | display
```

Figure 11: Output of `icat -o 10260 able2.dd 11108 | display`

2 Deleted File Identification and Recovery (Ext4)

In this exercise, you will be exposed to the limitations of the techniques introduced above when analyzing file systems where the metadata's block pointers are cleared upon deletion of a file; such is the case of Ext3 and Ext4 file systems. To demonstrate this feature, we will use the `able_3.dd` image as our target for analysis: an Ext4 file system. Before we begin, download the Ext4 disk image as follows:

```
root@kali:/volumes/storage1/Tut3/carve# wget http://turbina.gsd.inesc-id.pt  
/csf2021/able_3.tar.gz  
root@kali:/volumes/storage1/Tut3/carve# tar -xvzf able_3.tar.gz
```

This disk image is divided on a set of four split images. However, since TSK supports the inspection of individual split images, we may resort to TSK tools directly and do not need to fuse all splits of the image (something that can be accomplished by using `affuse`, for instance). Start by running `mmls able_3/able_3.000` for inspecting the first split (Figure 12). We are particularly interested in examining the `/home` directory, which is on the partition at offset 104448.

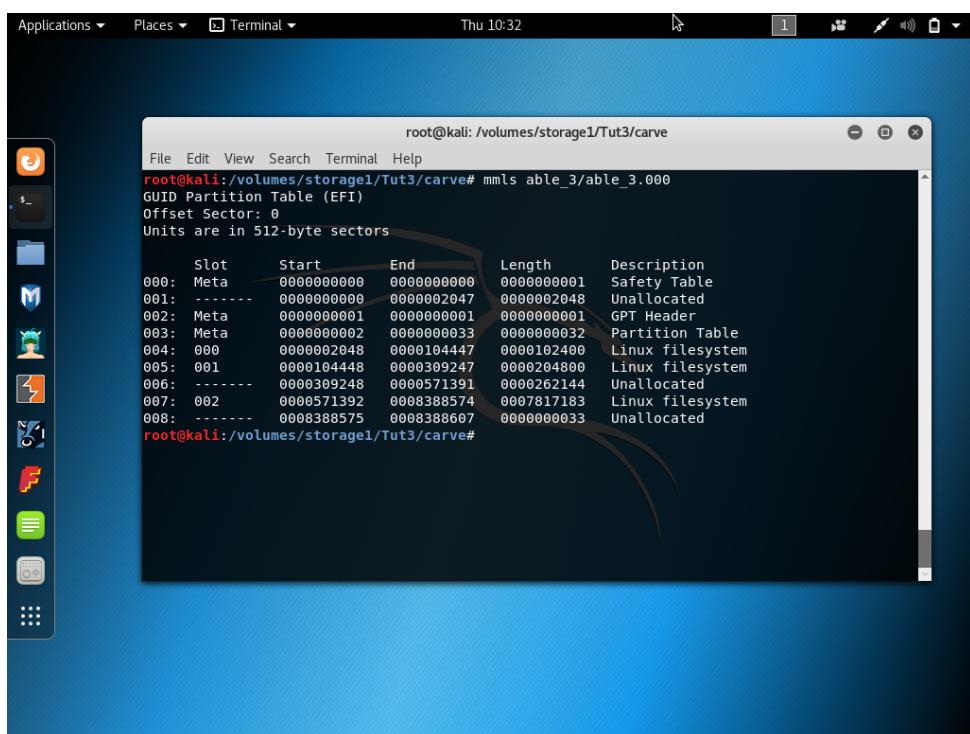


Figure 12: `mmls able_3/able_3.000`

You can quickly run `f1s -o 104448 -r able_3/able_3.000` to check which files are present in `/home` directory (Figure 13). You can see some familiar files in this output. We see a number of `lolitaz` files we saw on `able2`, and we also see the `lrkn.tar.gz` file from which we recovered the `README` from. For this exercise, we will be interested in the `lolitaz` files.

```

root@kali:~/volumes/storage1/Tut3/carve
File Edit View Search Terminal Help
root@kali:~/volumes/storage1/Tut3/carve# fls -o 104448 -r able_3/able_3.000
d/d 11: lost+found
d/d 12: ftp
d/d 13: albert
+ d/d 14: .h
++ r/d * 15(realloc): lolit_pics.tar.gz
++ r/r * 16(realloc): lolitaz1
++ r/r * 17: lolitaz10
++ r/r * 18: lolitaz11
++ r/r * 19: lolitaz12
++ r/r * 20: lolitaz13
++ r/r * 21: lolitaz2
++ r/r * 22: lolitaz3
++ r/r * 23: lolitaz4
++ r/r * 24: lolitaz5
++ r/r * 25: lolitaz6
++ r/r * 26: lolitaz7
++ r/r * 27: lolitaz8
++ r/r * 28: lolitaz9
+ d/d 15: Download
++ r/r 16: index.html
++ r/r * 17: lrkn.tar.gz
d/d 25689: $OrphanFiles
root@kali:~/volumes/storage1/Tut3/carve#

```

Figure 13: fls -o 104448 -r able_3/able_3.000

There is a single allocated file in that directory called `lolitaz13`. You can compare the output of `istat` and a follow-up `icat` command between the allocated file `lolitaz13` (inode 20), and one of the deleted files - we'll use `lolitaz2` (inode 21).

Figure 14 shows the output of `istat -o 104448 able_3.000 20`, where we observe that the inode is allocated and that the data it points to can be found in the direct blocks listed in the bottom.

```

root@kali:~/volumes/storage1/Tut3/carve
File Edit View Search Terminal Help
root@kali:~/volumes/storage1/Tut3/carve# istat -o 104448 able_3.000 20
inode: 20
Allocated
Group: 0
Generation Id: 1815721463
uid / gid: 1000 / 100
mode: rrw-r--r--
Flags: Extents,
size: 15045
num of links: 1

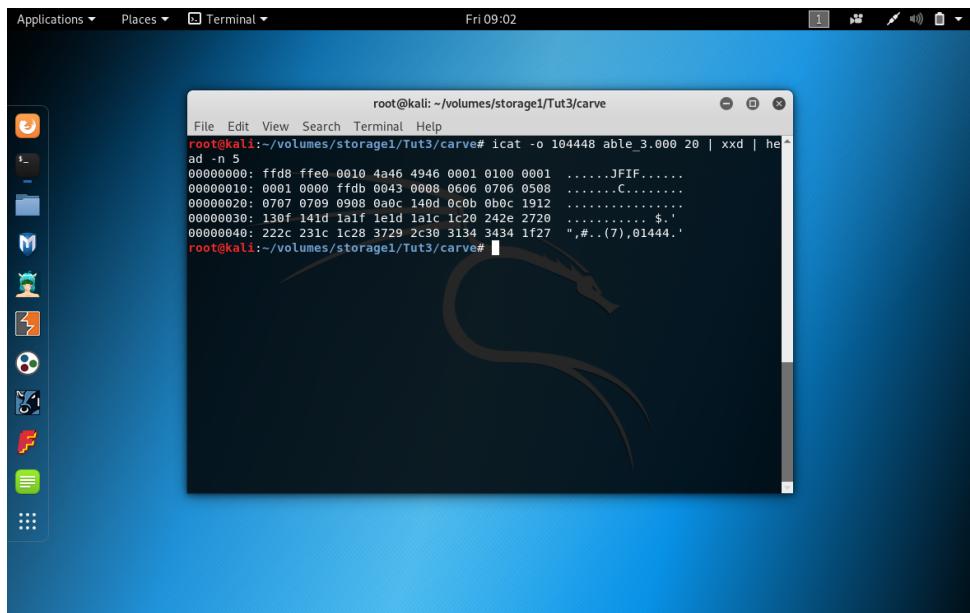
Inode Times:
Accessed: 2017-05-08 05:18:16 (BST)
File Modified: 2003-08-04 00:15:07 (BST)
Inode Modified: 2017-05-08 05:18:16 (BST)

Direct Blocks:
9921 9922 9923 9924 9925 9926 9927 9928
9929 9930 9931 9932 9933 9934 9935
root@kali:~/volumes/storage1/Tut3/carve#

```

Figure 14: Output of istat -o 104448 able_3.000 20

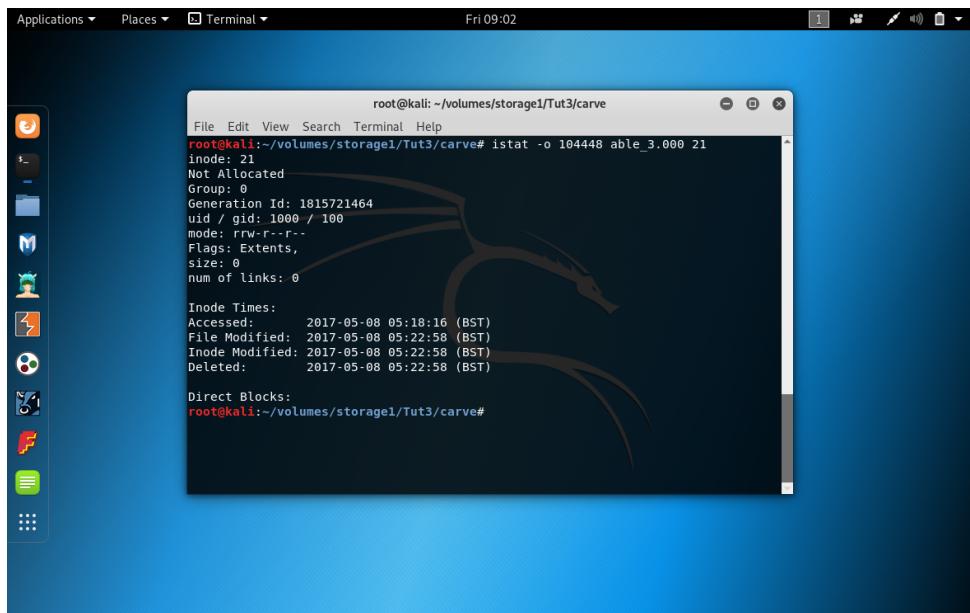
We can now inspect the initial content pertaining to this inode by running `icat -o 104448 able_3.000 20 | xxd | head -n 5`. The output of this command can be found in Figure 15 and shows the expected signature of a jpg image.



A screenshot of a Kali Linux desktop environment. A terminal window titled "root@kali: ~/volumes/storage1/Tut3/carve" is open, displaying the command "root@kali:~/volumes/storage1/Tut3/carve# icat -o 104448 able_3.000 20 | xxd | head -n 5". The output shows binary data starting with "ffd8 ffef 0010 4a46 4946 0001 0100 0001JFIF.....". The terminal window has a dark blue background with a Kali Linux logo watermark.

Figure 15: Output of `icat -o 104448 able_3.000 20 | xxd | head -n 5`

In contrast, we observe that inode 21 points to an unallocated file (Figure 16). As such, executing `icat -o 104448 able_3.000 21 | xxd | head -n 5` yields no output. On an Ext4 file system, when an inode is unallocated the entry for the Direct Blocks is cleared. There is no longer a pointer to the data, so commands like `icat` will not work. Remember that `icat` uses the information found in the inode to recover the file. In this case, there is none.



A screenshot of a Kali Linux desktop environment. A terminal window titled "root@kali: ~/volumes/storage1/Tut3/carve" is open, displaying the command "root@kali:~/volumes/storage1/Tut3/carve# istat -o 104448 able_3.000 21". The output shows the following details for inode 21:
inode: 21
Not Allocated
Group: 0
Generation Id: 1815721464
uid / gid: 1000 / 100
mode: rrw-r--r--
Flags: Extents,
size: 0
num of links: 0

Inode Times:
Accessed: 2017-05-08 05:18:16 (BST)
File Modified: 2017-05-08 05:22:58 (BST)
Inode Modified: 2017-05-08 05:22:58 (BST)
Deleted: 2017-05-08 05:22:58 (BST)

Direct Blocks:
root@kali:~/volumes/storage1/Tut3/carve#

Figure 16: Output of `istat -o 104448 able_3.000 21`

Albeit we are unable to recover the data of unallocated files through the above method, we can still apply a number of techniques to attempt the recovery of deleted files. The next section introduces the file carving technique which we will use to recover multiple `lolitaz` files.

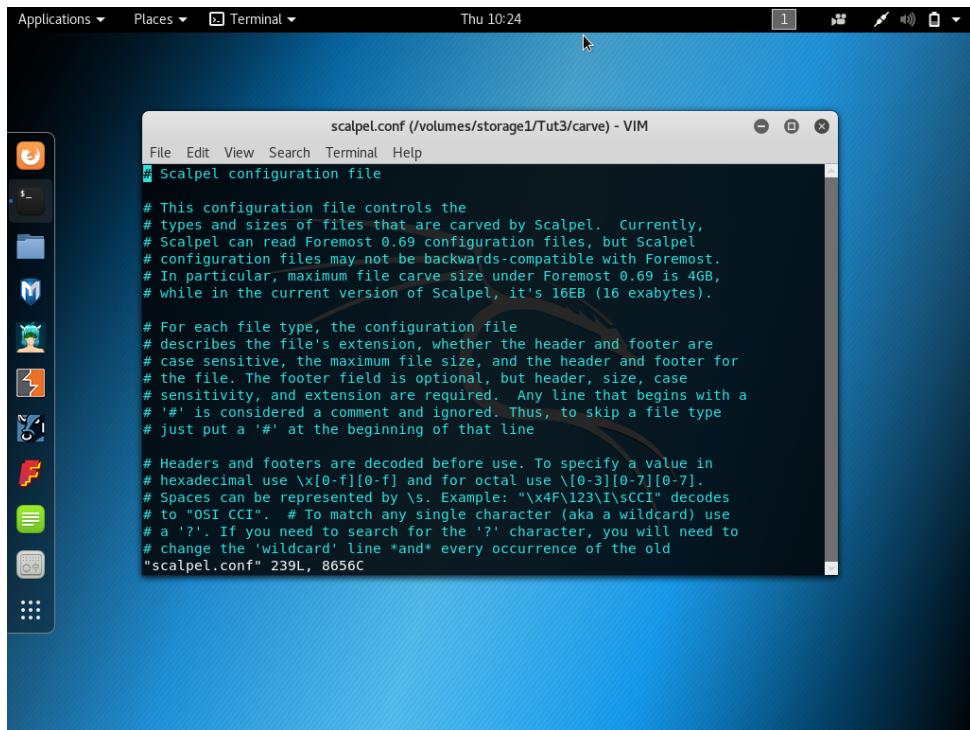
3 File Carving

File carving is a method for recovering files and fragments of files when directory entries are corrupt or missing. In a nutshell, file carving is a process used to extract structured data out of raw data present in a storage device, without the assistance of metadata provided by the file system that originally created the file. We will be using two well known carving tools, Scalpel (Section 3.1) and Foremost (Section 3.2), so as to retrieve files based on specific characteristics present in the structured data.

3.1 Scalpel

Scalpel is a filesystem-independent carver that reads a database of header and footer definitions and extracts matching files or data fragments from a set of image files or raw device files.

Before running Scalpel, you must define which file types are to be carved by the tool. Scalpel's configuration file (`/etc/scalpel/scalpel.conf`) starts out completely commented out (Figure 17). We will need to uncomment some file definitions in order to have Scalpel work.



```
scalpel.conf (/volumes/storage1/Tut3/carve) - VIM
File Edit View Search Terminal Help
# Scalpel configuration file

# This configuration file controls the
# types and sizes of files that are carved by Scalpel. Currently,
# Scalpel can read Foremost 0.69 configuration files, but Scalpel
# configuration files may not be backwards-compatible with Foremost.
# In particular, maximum file carve size under Foremost 0.69 is 4GB,
# while in the current version of Scalpel, it's 16EB (16 exabytes).

# For each file type, the configuration file
# describes the file's extension, whether the header and footer are
# case sensitive, the maximum file size, and the header and footer for
# the file. The footer field is optional, but header, size, case
# sensitivity, and extension are required. Any line that begins with a
# '#' is considered a comment and ignored. Thus, to skip a file type
# just put a '#' at the beginning of that line

# Headers and footers are decoded before use. To specify a value in
# hexadecimal use \x[0-f][0-f] and for octal use \[0-3][0-7][0-7].
# Spaces can be represented by \s. Example: "\x4F\123\1sCCI" decodes
# to "OSI CCI". # To match any single character (aka a wildcard) use
# a '?'. If you need to search for the '?' character, you will need to
# change the 'wildcard' line *and* every occurrence of the old

"scalpel.conf" 239L, 8656C
```

Figure 17: Scalpel config file

First, you should copy `/etc/scalpel/scalpel.conf` and edit it in your working directory (Figure 18), instructing Scalpel to perform an analysis according to this configuration file. Instead, you may also directly edit `/etc/scalpel/scalpel.conf`.

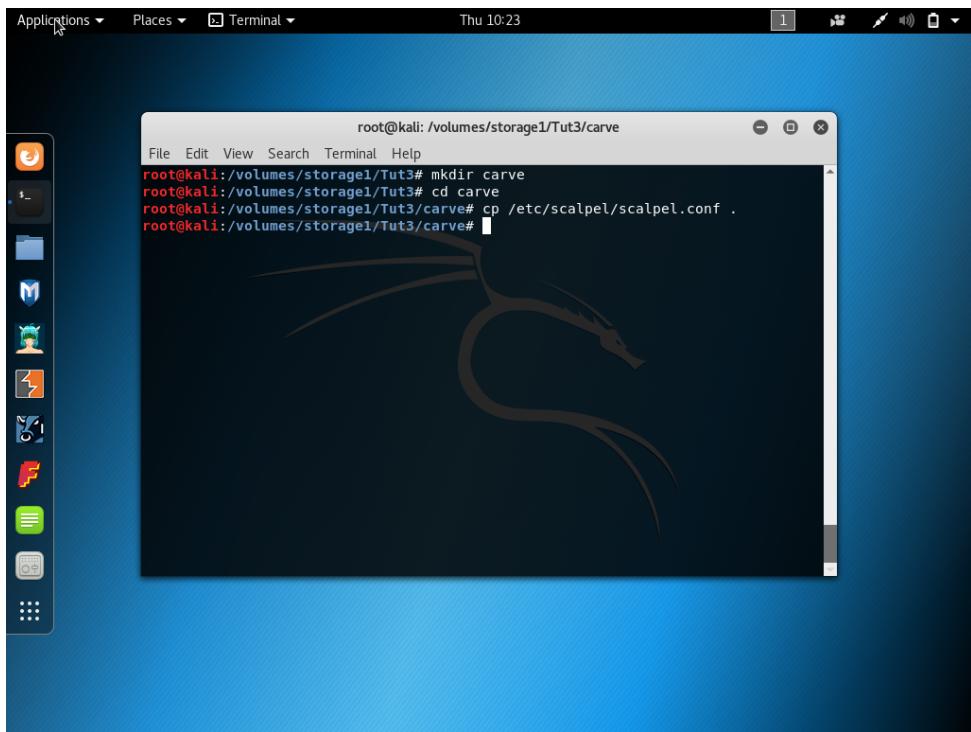


Figure 18: Setup environment

For the purpose of our exercise, scroll down to where the #GRAPHICS FILES section starts and uncomment every line that describes a file in that section. When we run Scalpel these uncommented lines will be used to search for patterns. That section should look like the one represented in Figure 19 when you are done.

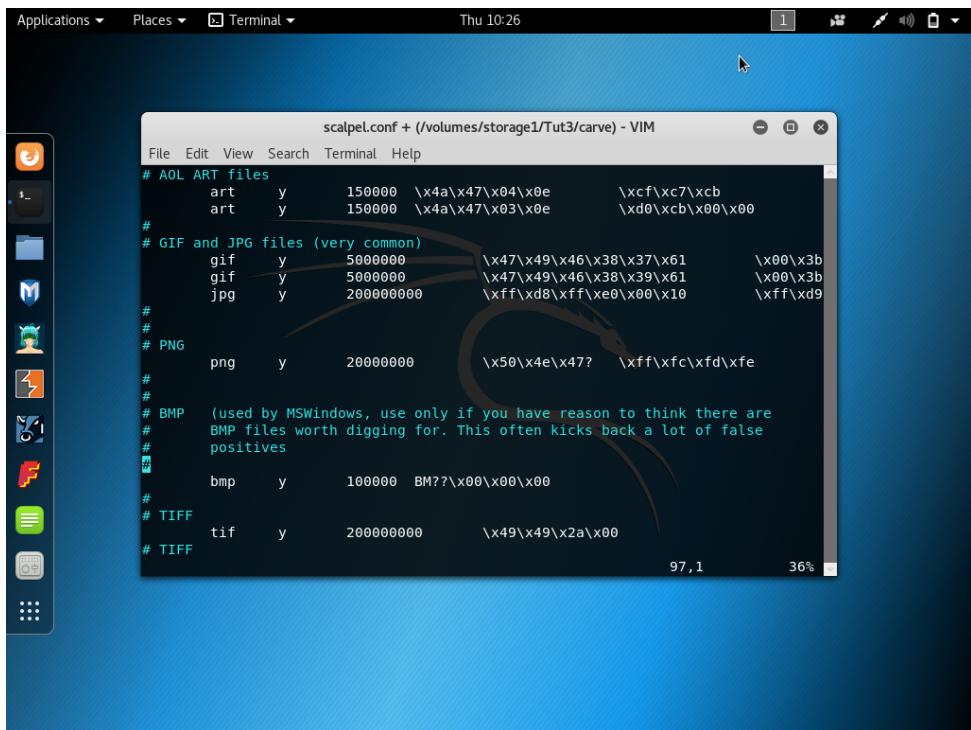


Figure 19: Scalpel config file changes

In this exercise, we aim at recovering the `lolitaz` files found in Section 2. Since we are able to retrieve allocated files with TSK tools, we shall focus our carving process on unallocated data only. We

can extract unallocated blocks by using the TSK tool `blkls`.

```
root@kali:/volumes/storage1/Tut3/carve# blkls -o 104448 able_3/able_3.000 > home.blkls
```

The `blkls` command is run with the offset (`-o`) pointing to the second Linux file system that starts at sector 104448. The output is redirected to `home.blkls`. The name “home” helps us to remember that this is the partition mounted as `/home`.

Scalpel has a number of options available to adjust the carving (be sure to check the manual!). There is an option to have Scalpel carve the files on block (or cluster) aligned boundaries. This means that you would be searching for files that start at the beginning of a data block. This should be done with caution. While you may get fewer false positives, it also means that you will miss files that may be embedded or “nested” in other files. Block aligned searching is done by providing the `-q <blocksize>` option. Try this option later, and compare the output. To get the block size for the target file system, you can use the `fsstat` command as we did in previous exercises.

In this case, we’ll use an option that allows us to specify which configuration file we want to use (`-c`). Finally, we’ll use the `-o` option to redirect our carved files to a directory we are going to call `scalpel_out` and the `-O` option so that the output remains in a single output directory instead of categorized sub directories. Having the files in a single folder makes it easier to inspect recovered evidence. With everything put together, you may run: `scalpel -c scalpel.conf -o scalpel_out -O home.blkls`. The output of the above command (Figure 20) shows scalpel carving those file types in which the definitions were uncommented.

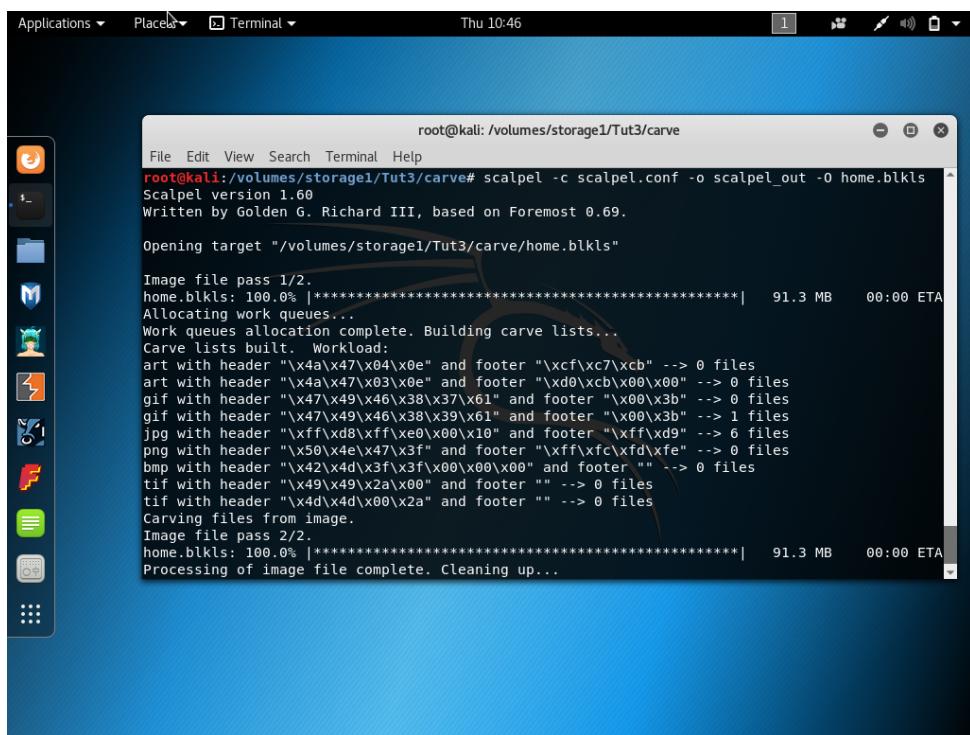


Figure 20: Scalpel usage

Once the command completes, a directory listing shows the carved files (Figure 21), and an `audit.txt` file providing a log with the contents of `scalpel.conf` and the program output. At the bottom of the output (Figure 22) is our list of carved files with the offset the header was found at, the length of the file, and the source (what was carved). The column labeled Chop would refer to files that had a maximum number of bytes carved before the footer was found.

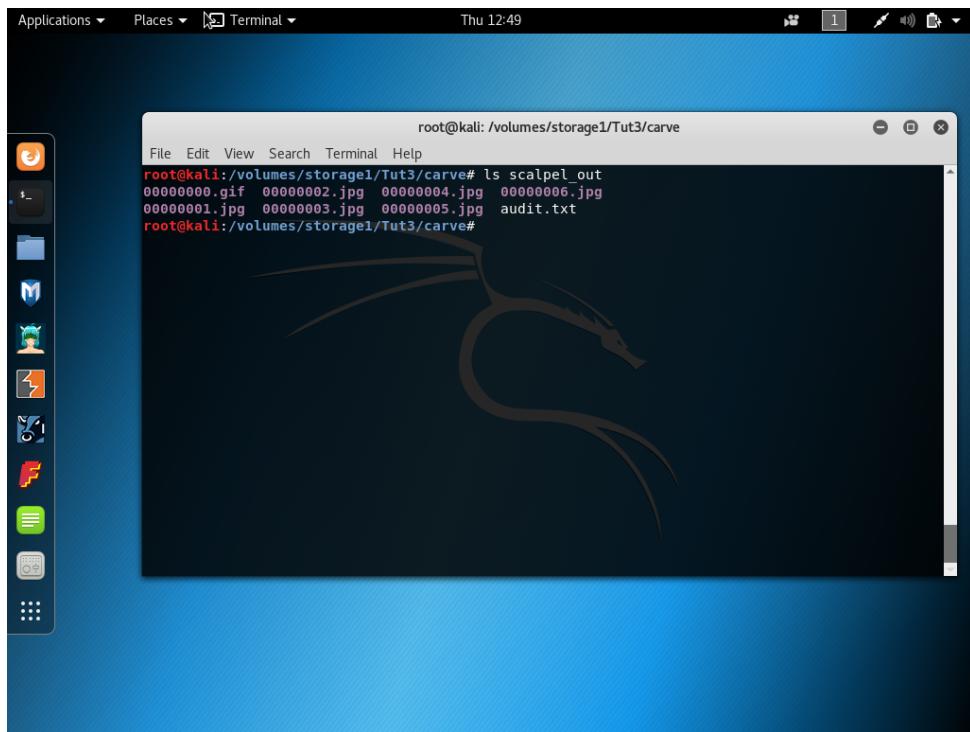


Figure 21: Scalpel output files

A screenshot of a terminal window titled "root@kali: /volumes/storage1/Tut3/carve". The window displays the Scalpel audit log. It starts with the version information "Scalpel version 1.60 audit file" and the start time "Started at Thu Oct 18 10:46:22 2018". The command used was "scalpel -c scalpel.conf -o scalpel_out -O home.blkls". The output directory is "/volumes/storage1/Tut3/carve/scalpel_out" and the configuration file is "scalpel.conf". It then shows the target directory being opened: "/volumes/storage1/Tut3/carve/home.blkls". The log then lists the files that were carved, showing their file names, start addresses, chop values, lengths, and the source block list (home.blkls). Finally, it shows the completion time "Completed at Thu Oct 18 10:46:25 2018" and the audit file path "scalpel_out/audit.txt (END)".

File	Start	Chop	Length	Extracted From
00000006.jpg	6586930	NO	6513	home.blkls
00000005.jpg	6586368	NO	7075	home.blkls
00000004.jpg	6278144	NO	15373	home.blkls
00000003.jpg	6249472	NO	27990	home.blkls
00000002.jpg	6129070	NO	5145	home.blkls
00000001.jpg	6128640	NO	5575	home.blkls
00000000.gif	6223872	NO	25279	home.blkls

Figure 22: Scalpel audit.txt

However, there are other files to be found in this unallocated data. To illustrate this, let's look at the `scalpel.conf` file again and add a different header definition for a bitmap file. Open `scalpel.conf` with your text editor and add the following line (highlighted in Figure 23) under the current `bmp` line in the `#GRAPHICS FILES` section.

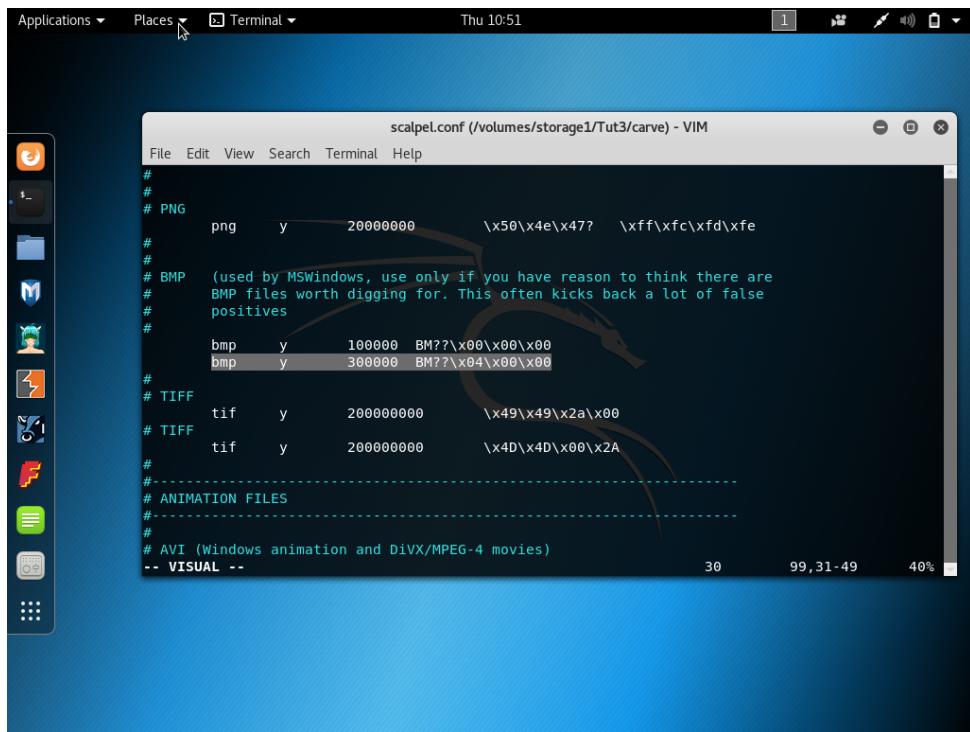


Figure 23: Scalpel config file changes

Here we changed the max size to 300000 bytes, and replaced the first `x00` string with `x04`. Save the file, re-run scalpel (write to a different output directory – `scalpel_out2`), and check the output (Figure 24).

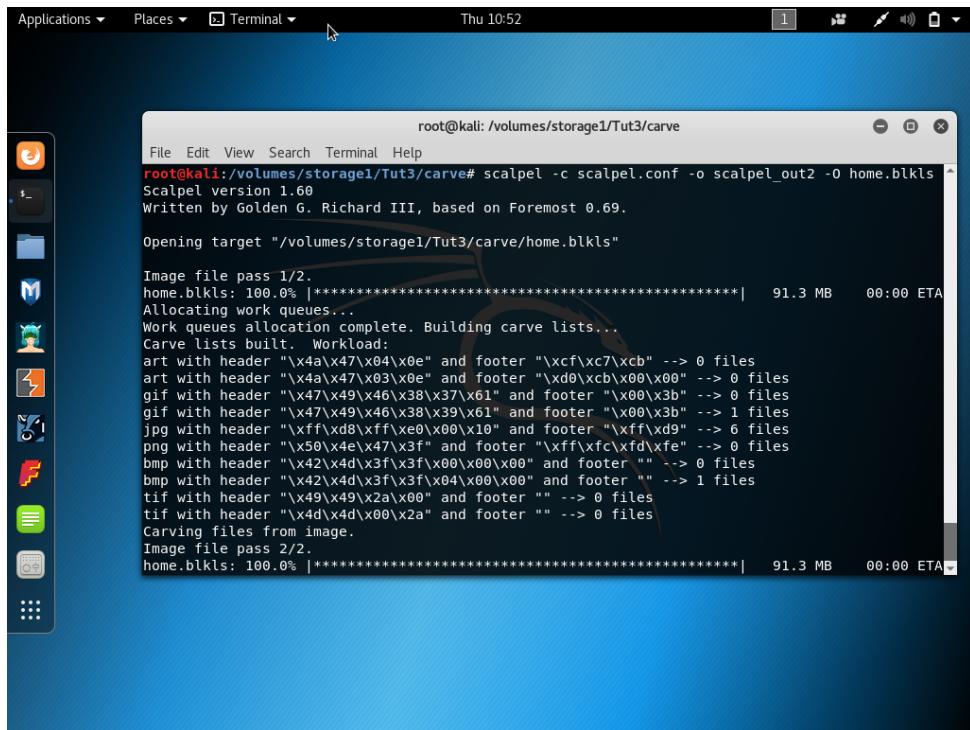


Figure 24: New scalpel output

Looking at the output above, we can see that a total of eight files were carved this time. The bitmap definition we added shows the `scalpel.conf` file can be easily improved on. Simply using `xxd` to find matching patterns in groups of files can be enough for you to build a decent library of headers, particularly if you come across many proprietary formats.

3.2 Foremost

File carving can be approached with a variety of tools, such as `foremost`, which also retrieves files by examining its internal structure. In this exercise, refer to `foremost` manual¹ and online walkthroughs² to carve out the files from `home.blks`. Does `foremost` yield the same results as Scalpel?

¹<https://www.systutorials.com/docs/linux/man/8-foremost/>

²<https://www.obsidianforensics.com/blog/deleted-file-recovery-using-foremost>